

Capstone Project



1. Problem Statement:

A **Retail Store** that has multiple outlets across the country.

They are facing issues in managing the inventory - to match the demand with respect to supply.

Dataset Information:

You are provided with the **Weekly_Sales** data for their various outlets.

The **Walmart.csv** contains 6435 rows and 8 columns.

Feature Name	Description
Stores	Store number
Date	Week of Sales
Weekly_Sales	Sales for the given store in that week
Holiday_Flag	Indicates if it is a holiday week
Temperature	Temperature on the day of the sale
Fuel_Price	Cost of fuel in the region

CPI

Consumer Price Index

2. Objectives:

1. Use

- ♦ Handle the missing values
- ♦ Exploratory Data Analysis (EDA)
- ♦ Outlier Analysis
- ♦ Statistical Analysis

To come up with various INSIGHTS that can give them a clear perspective on the following:

- ♦ a. If the Weekly_Sales are affected by the Unemployment Rate,
 - If YES - which Stores are suffering the most?
- ♦ b. If the Weekly_Sales show a seasonal trend,
 - When and what could be the reason?
- ♦ c. Does temperature affect the Weekly_Sales in any manner?
- ♦ d. How is the Consumer Price index (CPI) affecting the Weekly_Sales of various Stores?
- ♦ e. TOP Performing Stores according to the historical data.
- ♦ f. The WORST Performing Stores, and
 - How significant is the difference between the HIGHEST and LOWEST Performing Stores.

2. Use

- ♦ Predictive Modeling Techniques

To FORECAST the Weekly_Sales for each Stores for the NEXT 12 WEEKS.

Importing Libraries

```
In [8]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
import plotly.graph_objects as go
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PowerTransformer, FunctionTransformer
from sklearn.compose import ColumnTransformer
from sklearn.metrics import r2_score
from sklearn.tree import DecisionTreeRegressor
```

```

from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from statsmodels.stats.outliers_influence import variance_inflation_factor

```

```
In [9]: df_Walmart = pd.read_csv('Walmart.csv')
```

```
In [10]: df_Walmart
```

```
Out[10]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106
...
6430	45	28-09-2012	713173.95	0	64.88	3.997	192.013558	8.684
6431	45	05-10-2012	733455.07	0	64.89	3.985	192.170412	8.667
6432	45	12-10-2012	734464.36	0	54.47	4.000	192.327265	8.667
6433	45	19-10-2012	718125.53	0	56.47	3.969	192.330854	8.667
6434	45	26-10-2012	760281.43	0	58.85	3.882	192.308899	8.667

6435 rows × 8 columns

Basic EDA and Data Cleaning

Checking basic information of data

```
In [11]: df_Walmart.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Store           6435 non-null  int64  
1   Date            6435 non-null  object  
2   Weekly_Sales    6435 non-null  float64 
3   Holiday_Flag    6435 non-null  int64  
4   Temperature     6435 non-null  float64 
5   Fuel_Price      6435 non-null  float64 
6   CPI             6435 non-null  float64 
7   Unemployment    6435 non-null  float64 
dtypes: float64(5), int64(2), object(1)
memory usage: 402.3+ KB
```

Changing Data types

```
In [12]: #converting Date columns to data
df_Walmart['Date'] = pd.to_datetime(df_Walmart['Date'])
```

```
In [13]: #converting Holiday flag to bool
df_Walmart['Holiday_Flag'] = df_Walmart['Holiday_Flag'].astype('category')
```

```
In [14]: #converting Store to category
df_Walmart['Store'] = df_Walmart['Store'].astype('category')
```

```
In [15]: # adding year column to
df_Walmart['Year'] = df_Walmart['Date'].dt.year

# adding week column to
df_Walmart['Week'] = df_Walmart['Date'].dt.week

# adding month column to
df_Walmart['Month'] = df_Walmart['Date'].dt.month
df_Walmart.head(5)
```

```
Out[15]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment	Year	Week
0	1	2010-05-02	1643690.90	0	42.31	2.572	211.096358	8.106	2010	17
1	1	2010-12-02	1641957.44	1	38.51	2.548	211.242170	8.106	2010	48
2	1	2010-02-19	1611968.17	0	39.93	2.514	211.289143	8.106	2010	7
3	1	2010-02-26	1409727.59	0	46.63	2.561	211.319643	8.106	2010	8
4	1	2010-05-03	1554806.68	0	46.50	2.625	211.350143	8.106	2010	18

```
In [16]: df_Walmart.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
----  -
0   Store           6435 non-null   category
1   Date            6435 non-null   datetime64[ns]
2   Weekly_Sales    6435 non-null   float64
3   Holiday_Flag    6435 non-null   category
4   Temperature     6435 non-null   float64
5   Fuel_Price      6435 non-null   float64
6   CPI             6435 non-null   float64
7   Unemployment    6435 non-null   float64
8   Year            6435 non-null   int64
9   Week            6435 non-null   int64
10  Month           6435 non-null   int64
dtypes: category(2), datetime64[ns](1), float64(5), int64(3)
memory usage: 466.7 KB
```

Checking NULL values

```
In [17]: df_Walmart.isnull().sum()
```

```
Out[17]: Store          0
Date          0
Weekly_Sales  0
Holiday_Flag  0
Temperature   0
Fuel_Price    0
CPI           0
Unemployment  0
Year          0
Week          0
Month         0
dtype: int64
```

Reading the data for elementary Statistics

```
In [18]: df_Walmart.describe()
```


```
Out[18]:
```

	Weekly_Sales	Temperature	Fuel_Price	CPI	Unemployment	Year	Week	Mo
count	6.435000e+03	6435.000000	6435.000000	6435.000000	6435.000000	6435.000000	6435.000000	6435.000
mean	1.046965e+06	60.663782	3.358607	171.578394	7.999151	2010.965035	26.000000	6.475
std	5.643666e+05	18.444933	0.459020	39.356712	1.875885	0.797019	14.511794	3.321
min	2.099862e+05	-2.060000	2.472000	126.064000	3.879000	2010.000000	1.000000	1.000
25%	5.533501e+05	47.460000	2.933000	131.735000	6.891000	2010.000000	14.000000	4.000
50%	9.607460e+05	62.670000	3.445000	182.616521	7.874000	2011.000000	26.000000	6.000
75%	1.420159e+06	74.940000	3.735000	212.743293	8.622000	2012.000000	38.000000	9.000
max	3.818686e+06	100.140000	4.468000	227.232807	14.313000	2012.000000	52.000000	12.000

```
In [19]: #looking at correaltion between columns
```


```
plt.figure(figsize=(10,6))
heatmap_data = df_Walmart.corr()
mask = np.triu(np.ones_like(heatmap_data, dtype=bool))
sns.heatmap(heatmap_data,annot=True,cmap='Oranges',mask = mask)
plt.title("Correlation between Columns",fontdict={'fontsize':20,'color':'Green','fontwei
plt.show()
```

Correlation between Columns



Outliers Detection

```
In [20]: # Using box plots
plt.figure(figsize=(18,8))
colours = sns.color_palette(n_colors=8)
sns.boxplot(data = df_Walmart,palette=colours)
plt.show()
```



Outliers Treatment

```
In [21]: Q3 = df_Walmart['Weekly_Sales'].quantile(0.75)
Q1 = df_Walmart['Weekly_Sales'].quantile(0.25)
IQR = Q3-Q1
upperLimit = Q3+(1.5*IQR)
lowerLimit = Q1-(1.5*IQR)
```

```
In [22]: filt = ((df_Walmart['Weekly_Sales']<= upperLimit) & (df_Walmart['Weekly_Sales']>= lowerLimit))
df_Walmart = df_Walmart[filt]
```

```
In [23]: df_Walmart
```

```
Out[23]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment	Year	Week
0	1	2010-05-02	1643690.90	0	42.31	2.572	211.096358	8.106	2010	1
1	1	2010-12-02	1641957.44	1	38.51	2.548	211.242170	8.106	2010	4
2	1	2010-02-19	1611968.17	0	39.93	2.514	211.289143	8.106	2010	
3	1	2010-02-26	1409727.59	0	46.63	2.561	211.319643	8.106	2010	
4	1	2010-05-03	1554806.68	0	46.50	2.625	211.350143	8.106	2010	1
...
6430	45	2012-09-28	713173.95	0	64.88	3.997	192.013558	8.684	2012	3
6431	45	2012-05-10	733455.07	0	64.89	3.985	192.170412	8.667	2012	1
6432	45	2012-12-10	734464.36	0	54.47	4.000	192.327265	8.667	2012	5
6433	45	2012-10-19	718125.53	0	56.47	3.969	192.330854	8.667	2012	4
6434	45	2012-10-26	760281.43	0	58.85	3.882	192.308899	8.667	2012	4

6401 rows × 11 columns

Checking for duplicate values

```
In [24]: df_Walmart.duplicated().sum()
```

```
Out[24]: 0
```

Insights that will give a clear perspective on the following

(a).If the Weekly_Sales are affected by the Unemployment Rate

Correlation between Weekly_Sales and Unemployment Rate

```
In [25]: #checking correaltion between Weekly Sales and Unemployment Rate
correlationCoeff = df_Walmart['Weekly_Sales'].corr(df_Walmart['Unemployment'])
print("The correlation coefficient of Weekly Sales and Unemployment Rate is : ",correlationCoeff)

The correlation coefficient of Weekly Sales and Unemployment Rate is : -0.10429750912578391
```


```
In [26]: #plotting the regression plot

plt.figure(figsize=(20,10))
sns.regplot(data = df_Walmart,x = 'Unemployment',y = 'Weekly_Sales')

# Set the title
plt.title('Unemployment vs Weekly Sales ',fontdict={'fontsize':20,'color':'Green','fontw

# Set the x and y axis labels
plt.xlabel('Unemployment', color='#FF8C00', fontweight='bold', fontsize=16)
plt.ylabel('Weekly Sales', color='#FF8C00', fontweight='bold', fontsize=16)

plt.show()
```



(b).Which Stores are suffering the most?

```
In [ ]:
```

```
In [27]: # getting the correaltion between weekly sales and unemployment and sorting it
sortwise_correlation = df_Walmart.groupby('Store')['Weekly_Sales', 'Unemployment'].corr

#restting the index
sortwise_correlation.reset_index(inplace = True)

#creating filter
```


```
min_corr = sortwise_correlation[('Weekly_Sales', 'Unemployment')].min()
filt = (sortwise_correlation[('Weekly_Sales', 'Unemployment')] == min_corr)
store = sortwise_correlation.loc[filt, 'Store']
print("Stores with the Highest Negative Correlation with Unemployment Rate:", store[0])
```

Stores with the Highest Negative Correlation with Unemployment Rate: 38

(c).If the Weekly_Sales show a seasonal trend?


When and what could be the reason?

```
In [28]: # checking for holiday flag and how it affects weekly sales
plt.figure(figsize=(14,8))
sns.set_style('darkgrid')
sns.boxplot(data = df_Walmart, x = 'Holiday_Flag', y = 'Weekly_Sales')
plt.title("Holidays vs Weekly Sales", fontdict={'fontsize':20, 'color':'Green', 'fontweight':
plt.xlabel('Holiday Flag', color='#FF8C00', fontweight='bold', fontsize=14)
plt.ylabel('Weekly Sales', color='#FF8C00', fontweight='bold', fontsize=14)
plt.show()
```



(d).Does temperature affect the Weekly_Sales in any manner?

```
In [29]: # checking for temperature and how it affects weekly sales with holiday flag
sns.set_style('darkgrid')
sns.lmplot(data = df_Walmart, x = 'Temperature', y = 'Weekly_Sales', hue='Holiday_Flag', hei
plt.title("Temperature vs Weekly Sales with Holidays", fontdict={'fontsize':20, 'color':'G
plt.xlabel('Temperature', color='#FF8C00', fontweight='bold', fontsize=16)
plt.ylabel('Weekly Sales', color='#FF8C00', fontweight='bold', fontsize=16)
plt.show()
```




(e).Changes in Temperature over time

```
In [30]: sns.set_style('darkgrid')
sns.relplot(data = df_Walmart, x = 'Month', y = 'Temperature', hue='Year',
            kind = 'line', height = 4, aspect=2.5, palette='Set2',
            marker = '^')

# Set the title
plt.title('Changes in Temperature over Time', fontdict={'fontsize':22,'color':'Maroon',})

# Set the x and y axis labels
plt.xlabel('Month', color='brown',fontweight='bold', fontsize=16)
plt.ylabel('Temperature', color='brown',fontweight='bold', fontsize=16)
plt.show()
```



(f).How is the Consumer Price Index (CPI) affecting the Weekly_Sales of various Stores?

```
In [31]: # Create a list of unique store numbers
stores = df_Walmart['Store'].unique()

# Set up subplots

fig, axes = plt.subplots(nrows=len(stores)//3, ncols=3, figsize=(15, 5*len(stores)//3))
fig.tight_layout(pad=5.0)

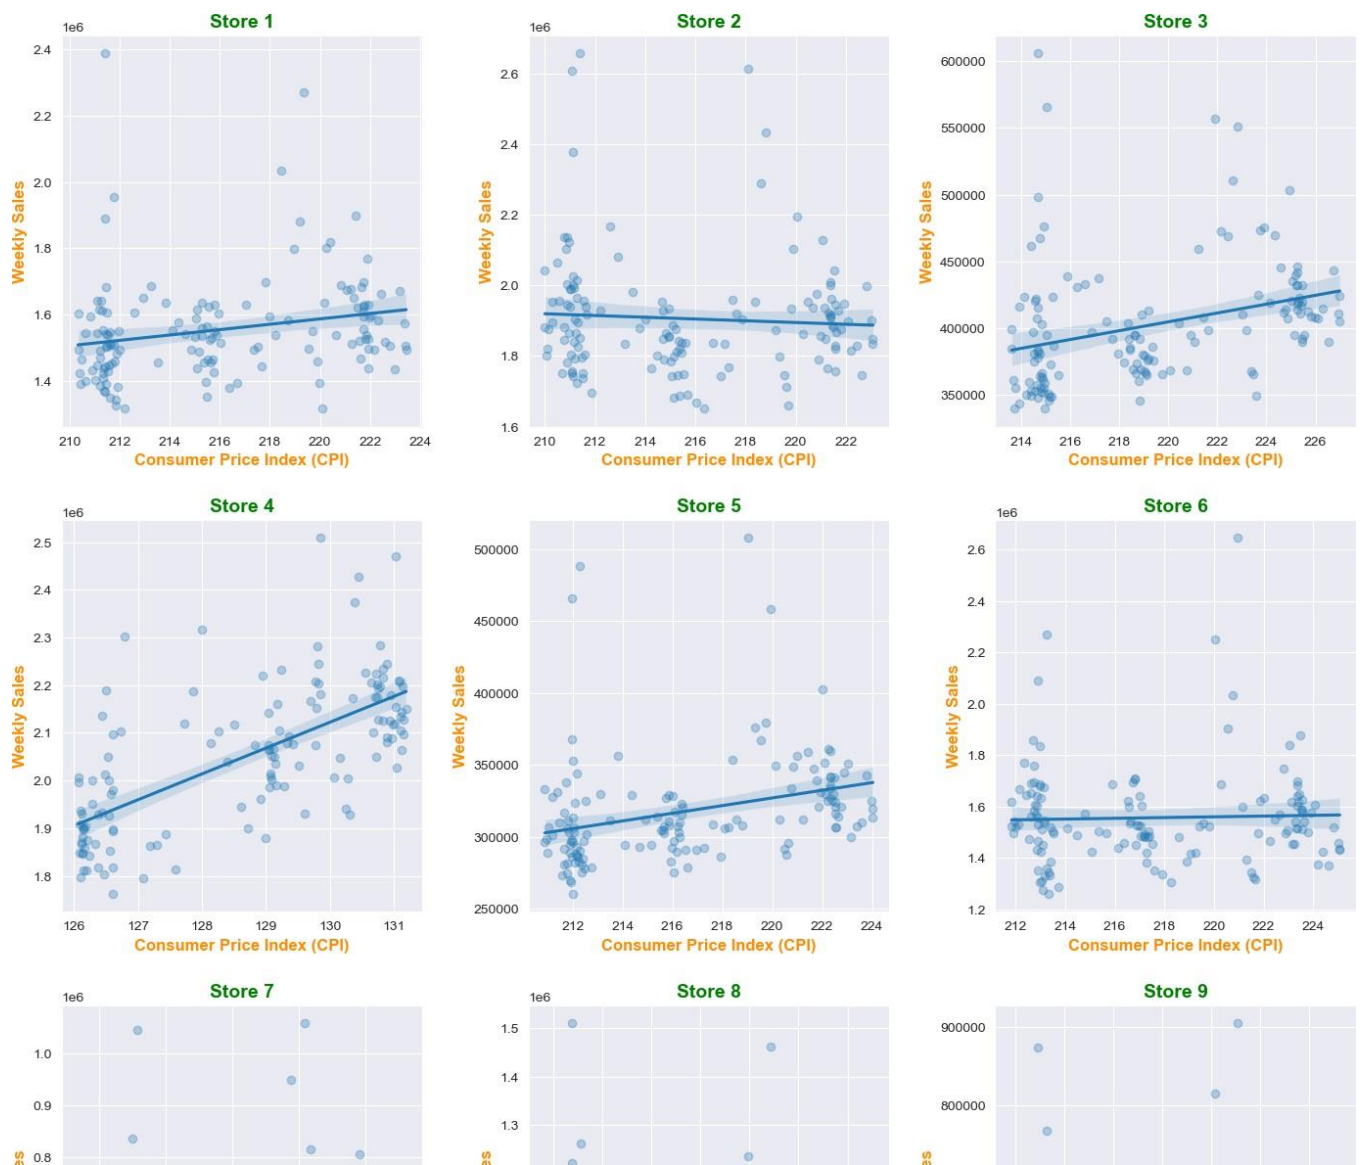
# Iterate over each store and create a scatter plot
for i, store_num in enumerate(stores):
    row = i // 3
    col = i % 3
    ax = axes[row, col]




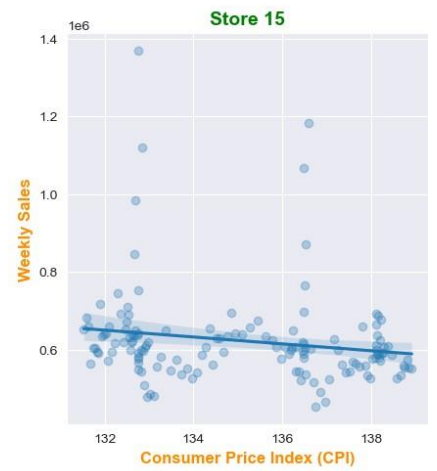


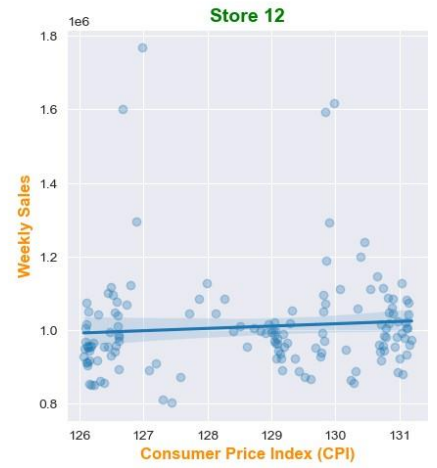


    # Filter the data for the specific store
    store_data = df_Walmart[df_Walmart['Store'] == store_num]










    # Plot scatter plot with regression line
    sns.set_style('darkgrid')
    sns.regplot(x='CPI', y='Weekly_Sales', data=store_data, ax=ax, scatter_kws={'alpha':

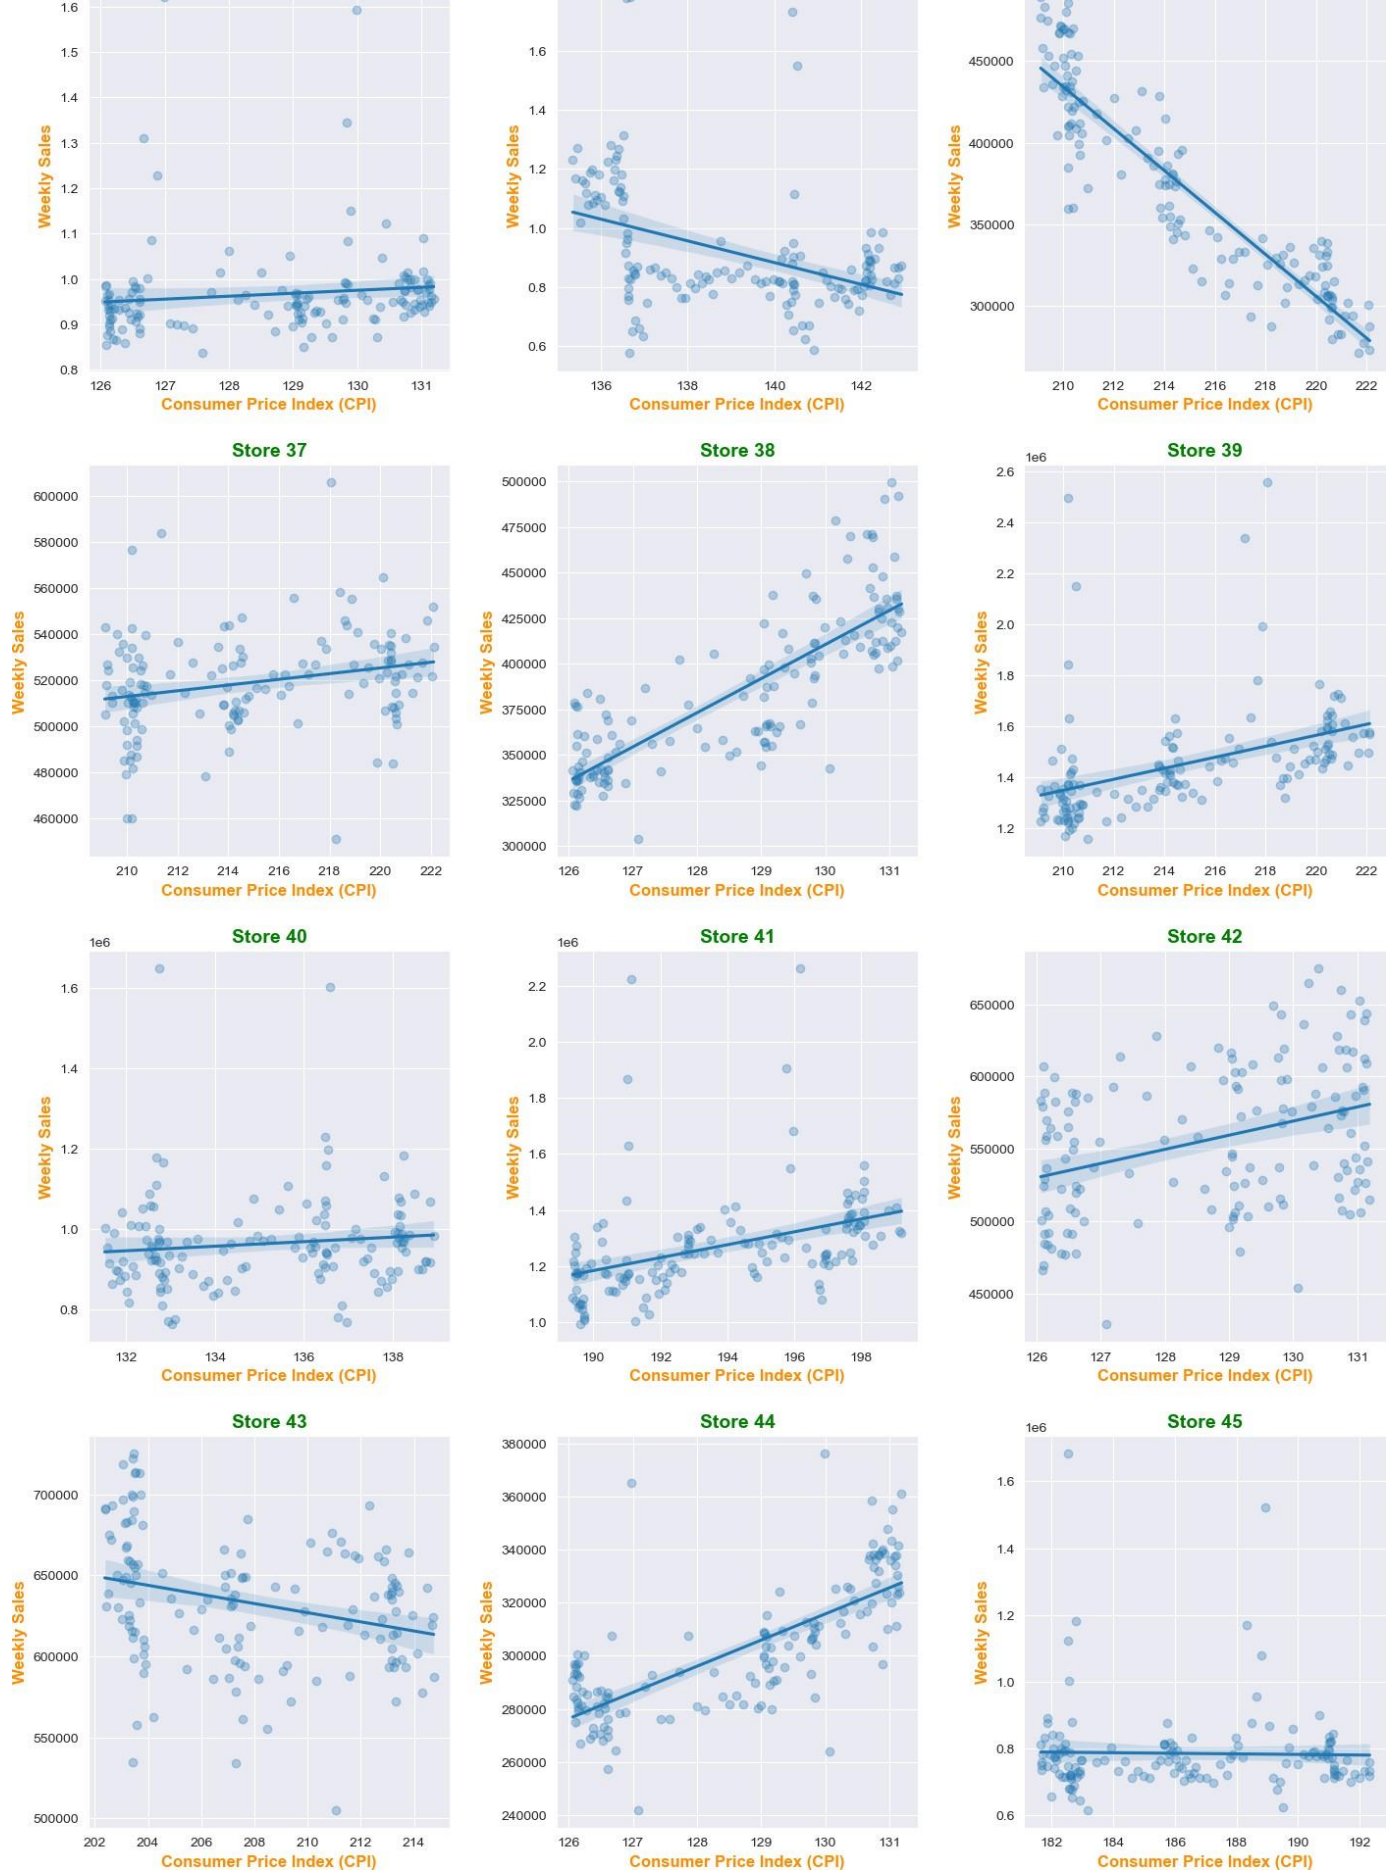
ax.set_title(f'Store {store_num}', fontdict={'fontsize':14, 'color':'Green', 'fontweigh
ax.set_xlabel('Consumer Price Index (CPI)', color='#FF8C00', fontweight='bold', fon
ax.set_ylabel('Weekly Sales', color='#FF8C00', fontweight='bold', fontsize=12)

plt.show()
```









(g).Changes in CPI (CONSUMER PRICE INDEX) over time

```
In [32]: plt.figure(figsize=(16,8))
sns.set_style('darkgrid')
```


```

sns.lineplot(
    data=df_Walmart,
    x='Month',
    y='Temperature',
    hue='Year',
    ci='sd', # Specify the standard deviation for shading
    err_style='band', # Use a shaded band for the confidence interval
    palette='Set1'
)

# Set the title
plt.title('Changes in CPI (CONSUMER PRICE INDEX) over Time', fontdict={'fontsize':22,'co

# Set the x and y axis labels
plt.xlabel('Month', color='orange',fontweight='bold', fontsize=16)
plt.ylabel('CONSUMER PRICE INDEX', color='orange',fontweight='bold', fontsize=16)
plt.show()

```



(h).How is the Fuel Prices are affecting the Weekly_Sales of various Stores?

```


In [33]: sns.set_style('darkgrid')
plt.figure(figsize=(20,8))
sns.regplot(data = df_Walmart,x = 'Fuel_Price',y = 'Weekly_Sales')

# Set the title
plt.title('Fuel Price vs Weekly Sales', fontdict={'fontsize':22,'color':'Green','fontwei

# Set the x and y axis labels
plt.ylabel('Weekly Sales', color='#FF8C00', fontweight='bold', fontsize=16)
plt.xlabel('Fuel Price', color='#FF8C00', fontweight='bold', fontsize=16)

plt.show()

```





```
In [34]: sns.set_style('darkgrid')
plt.figure(figsize=(12,6))
sns.barplot(data = df_Walmart,x = 'Year',y = 'Fuel_Price')

# Set the title
plt.title('Average Fuel Prices Over the Years', fontdict={'fontsize':22,'color':'Green',

# Set the x and y axis labels
plt.xlabel('Year', color='#FF8C00',fontweight='bold', fontsize=16)
plt.ylabel('Fuel Price', color='#FF8C00',fontweight='bold', fontsize=16)

plt.show()
```



(i).Relation between Date and Weekly Sales

```
In [35]: sns.set_style('darkgrid')
sns.relplot(data = df_Walmart,x = 'Date',y = 'Weekly_Sales',hue='Year',kind = 'line',hei
```

```
# Set the title
plt.title('Weekly Sales Over Time by Year', fontdict={'fontsize':22,'color':'Green','fon

# Set the x and y axis labels
plt.xlabel('Date', color='#FF8C00',fontweight='bold', fontsize=16)
plt.ylabel('Sales', color='#FF8C00',fontweight='bold', fontsize=16)
plt.show()
```




(j).Relation between Store and Weekly Sales

```
In [36]: plt.figure(figsize=(20,6))
sns.set_style('whitegrid')
sns.barplot(data = df_Walmart,x = 'Store',y = 'Weekly_Sales',hue = 'Year',palette='magma')

# Set the title
plt.title('Weekly Sales by Store and Year', fontdict={'fontsize':22,'color':'Green','fon

# Set the x and y axis labels
plt.xlabel('STORES', color='brown',fontweight='bold', fontsize=16)
plt.ylabel('WEEKLY SALES', color='brown',fontweight='bold', fontsize=16)
plt.xticks(fontweight='bold', fontsize=12)
plt.yticks(fontweight='bold', fontsize=12)
plt.show()
```



TOP Performing Stores according to the historical data

TOP Performing Stores Year wise according to the historical data

```

In [37]: #grouping the data Year wise and then Store wise
performers = df_Walmart.groupby(['Year', 'Store']).agg(Yearly_Sales=('Weekly_Sales', 'sum')


best_store = pd.DataFrame()
# getting unique values in Year column
year = df_Walmart.Year.unique()
for i in year:
    # Filter the top performers for the current year
    filt = (performers['Year'] == i)
    max_sale = performers[filt]['Yearly_Sales'].max()
    mask = (performers['Yearly_Sales'] == max_sale)
    best_store = pd.concat([best_store, performers[filt][mask]])

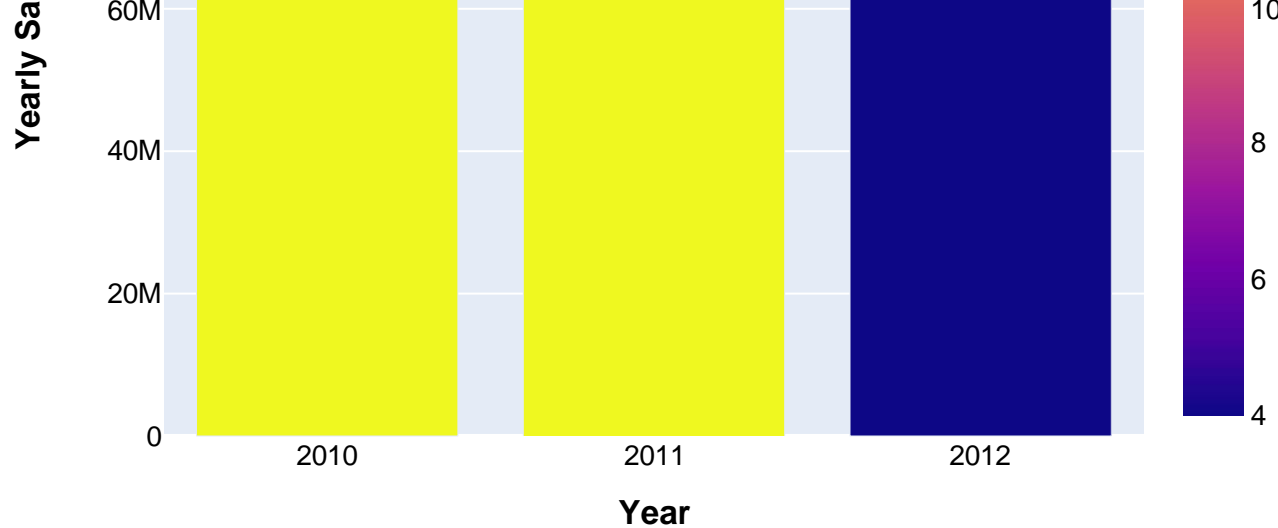
best_store.reset_index(drop = True, inplace=True)

#Changing the datatype to int for Store column
best_store['Store'] = best_store['Store'].astype('int')
unique_years = sorted(best_store['Year'].unique())
#plotting the graph for top performing stores
fig = px.bar(
    best_store,
    x='Year',
    y='Yearly_Sales',
    color='Store',
    color_discrete_sequence=px.colors.qualitative.Set1,
    category_orders={"Year": sorted(best_store['Year'].unique())}, # Enforce order with
)
fig.update_layout(
    xaxis=dict(
        tickmode='array',
        tickvals=unique_years,
        dtick=1, # Set the tick interval to 1 to display only integer values
    ),
    title='<b>Yearly Sales Trends: Top Performing Stores (2010-2012)</b>',
    xaxis_title='<b>Year</b>',
    yaxis_title='<b>Yearly Sales</b>',
    showlegend=True,
    font=dict(
        family="Arial, sans-serif",
        size=14,
        color="black"
    ),
    title_font=dict(
        family="Arial, sans-serif",
        size=22,
        color="green",
    )
)
fig.show()

```

Yearly Sales Trends: Top Performing Stores (2010-2012)





WORST Performing Stores according to the historical data

```
In [38]: worst_store = pd.DataFrame()
# getting unique values in Year column
year = df_Walmart.Year.unique()
for i in year:
    # Filter the top performers for the current year
    filt = (performers['Year']==i)
    min_sale = performers[filt]['Yearly_Sales'].min()
    mask = (performers[filt]['Yearly_Sales']==min_sale)
    worst_store = pd.concat([worst_store,performers[filt][mask]])

worst_store.reset_index(drop = True,inplace=True)


#Changing the datatype to int for Store column
worst_store['Store'] = worst_store['Store'].astype('int')
unique_years = sorted(worst_store['Year'].unique())
#plotting the graph for top performing stores
fig = px.bar(
    worst_store,
    x='Year',
    y='Yearly_Sales',
    color='Store',
    color_discrete_sequence=px.colors.qualitative.Set1,
    category_orders={"Year": sorted(worst_store['Year'].unique())}, # Enforce order wit
)
fig.update_layout(
    xaxis=dict(
        tickmode='array',
        tickvals=unique_years,
        dtick=1, # Set the tick interval to 1 to display only integer values
    ),
    title='<b>Yearly Sales Trends: Worst Performing Stores (2010-2012)</b>',
    xaxis_title='<b>Year</b>',
    yaxis_title='<b>Yearly Sales</b>',
    showlegend=True,
    font=dict(
        family="Arial, sans-serif",
        size=14,
        color="black"
    ),
    title_font=dict(
        family="Arial, sans-serif",
```

```

        size=22,
        color="green",
    )
)
fig.show()

```

Yearly Sales Trends: Worst Performing Stores (2010-2012)



How significant is the difference between the HIGHEST and LOWEST Performing Stores

In [39]: *# To assess the significance of the difference between the highest and lowest performing stores we will calculate various statistical measures.*

```

for i in year:
    filt = (performers['Year']==i)
    sales_range = performers[filt]['Yearly_Sales'].max() - performers[filt]['Yearly_Sale
    print('Stats for Yearly Sales in the year ',i)
    print('Range : ',sales_range)

    sales_mean = performers[filt]['Yearly_Sales'].mean() # Avg sales
    print('Average Sales : ',sales_mean)

    sales_std = performers[filt]['Yearly_Sales'].std() #std deviation
    print('Standtard Deviation : ',sales_std)

    cv = (sales_std / sales_mean) * 100 #coefficient of Variance
    print('Relative Variance : ',cv)

    #percentile analysis
    sales_25th_percentile = performers[filt]['Yearly_Sales'].quantile(0.25)
    sales_75th_percentile = performers[filt]['Yearly_Sales'].quantile(0.75)
    print('25th_percentile : ',sales_25th_percentile)

```

```
print('75th_percentile : ',sales_75th_percentile)
print(' ')
```

Stats for Yearly Sales in the year 2010
Range : 83192150.55
Average Sales : 49501474.735999994
Standtard Deviation : 24543508.62317917
Relative Variance : 49.58136854320803
25th_percentile : 25568078.15
75th_percentile : 65782276.32


Stats for Yearly Sales in the year 2011
Range : 89769365.03999999
Average Sales : 53434914.90044444
Standtard Deviation : 26897017.527062986
Relative Variance : 50.33603511332488
25th_percentile : 29117302.669999998
75th_percentile : 74169225.52

Stats for Yearly Sales in the year 2012
Range : 81335638.21
Average Sales : 44447396.87444445
Standtard Deviation : 23019092.759410933
Relative Variance : 51.78951834780234
25th_percentile : 24827530.71
75th_percentile : 59212433.28

Distribution graph of columns


Temperature

```
In [40]: plt.figure(figsize=(10,6))
sns.distplot(df_Walmart['Temperature'])
plt.show()
```




Fuel Price

```
In [41]: plt.figure(figsize=(10,6))  
sns.distplot(df_Walmart['Fuel_Price'])  
plt.show()
```




CPI

```
In [42]: plt.figure(figsize=(10,6))
sns.distplot(df_Walmart['CPI'])
plt.show()
```



Unemployment

```
In [43]: plt.figure(figsize=(10,6))
sns.distplot(df_Walmart['Unemployment'])
plt.show()
```




```
In [44]: df_Walmart.skew()
```

```
Out[44]: Weekly_Sales    0.497428
Temperature -0.344571
Fuel_Price  -0.101046
CPI          0.061319
Unemployment 1.186512
Year         0.057389
Week         0.021496
Month        0.027883
dtype: float64
```

VIF check for multicollinearity

```
In [45]: lst = list(df_Walmart.columns)
lst.remove('Weekly_Sales')
lst.remove('Store')
lst.remove('Holiday_Flag')
lst.remove('Date')
X = df_Walmart[lst]
y = df_Walmart['Weekly_Sales']
```

```
In [46]: # VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif_data)
```

	feature	VIF
0	Temperature	13.291808
1	Fuel_Price	60.255810
2	CPI	24.527554
3	Unemployment	22.061234
4	Year	145.332635
5	Week	606.455295
6	Month	691.535853

```
In [47]: #we can drop year,month,week columns as those columns were not there in the original dat
df_Walmart.drop(columns=['Year','Week','Month'],inplace = True)
```

Insights from distribution and skewness for selection of transformers

Temperature & Fuel need Power

Transformer Unemployment needs

Function Transformer

To select features we need to do a statistical assessment to remove multicollinear columns

```
In [48]: X = df_Walmart[['Temperature', 'Fuel_Price', 'CPI', 'Unemployment']]
y = df_Walmart['Weekly_Sales']

model = sm.OLS(y, sm.add_constant(X)).fit()
```

```
In [49]: print(model.summary())
```

OLS Regression Results

=====

D
e
p
.

V
a
r
i
a
b
l
e
:

W
e
e
k
l
y

-
S
a
l
e
s

R
-
s
q
u
a
r
e
d
:

0
.
0
2
2

```

Model: OLS Adj. R-squared: 0.022
Method: Least Squares F-statistic: 36.53
Date: Sun, 26 Nov 2023 Prob (F-statistic): 3.06e-30
Time: 22:17:14 Log-Likelihood: -93560.
No. Observations: 6401 AIC: 1.871e+05
Df Residuals: 6396 BIC: 1.872e+05
Df Model: 4
Covariance Type: nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.651e+06	7.72e+04	21.378	0.000	1.5e+06	1.8e+06
Temperature	-318.2821	384.464	-0.828	0.408	-1071.960	435.396
Fuel_Price	-4817.1235	1.53e+04	-0.316	0.752	-3.47e+04	2.51e+04
CPI	-1524.2365	189.384	-8.048	0.000	-1895.493	-1152.980
Unemployment	-3.971e+04	3848.319	-10.319	0.000	-4.73e+04	-3.22e+04
Omnibus:		555.168	Durbin-Watson:		0.090	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		370.723	
Skew:		0.474	Prob(JB):		3.15e-81	
Kurtosis:		2.299	Cond. No.		2.16e+03	

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.16e+03. This might indicate that there are strong multicollinearity or other numerical problems.

We will select models which are robust to multicollinearity or handle it gracefully

Decision Tree

Random Forest

XGBoost

Predictive Models:

FEATURE TRANSFORMATION AND SELECTION:

Feature Transformation and Predicting

```

In [50]: #setting Date column to index
df_Walmart.set_index('Date', inplace = True)

```

```

In [51]: df_Walmart.head()

```

```

Out[51]:
      Store  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price      CPI  Unemployment
Date
2010-05-02    1    1643690.90           0         42.31         2.572  211.096358         8.106
2010-12-02    1    1641957.44           1         38.51         2.548  211.242170         8.106

```

2010-02-19	1	1611968.17	0	39.93	2.514	211.289143	8.106
2010-02-26	1	1409727.59	0	46.63	2.561	211.319643	8.106
2010-05-03	1	1554806.68	0	46.50	2.625	211.350143	8.106

```
In [52]: lst = list(df_Walmart.columns)
lst.remove('Weekly_Sales')
X = df_Walmart[lst]
y = df_Walmart['Weekly_Sales']
```

```
In [53]: function_list = ['Unemployment']
power_list = ['Temperature', 'Fuel_Price']
skip_list = ['Store', 'Holiday_Flag', 'CPI']
```

```
In [54]: transformers = [('function', FunctionTransformer(np.log1p), function_list),
                        ('power', PowerTransformer(), power_list),
                        ('skip', 'passthrough', skip_list)]
```

```
In [55]: column_Transformer = ColumnTransformer(transformers=transformers, remainder = 'passthrough')
```

```
In [56]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)

results_df = pd.DataFrame(columns=['Model', 'y_test', 'y_pred', 'R2 Score'])
# Define models and their respective hyperparameter grids
models = {
    'Decision Tree': (DecisionTreeRegressor(), {'model__max_depth': [None, 5, 10, 15]}),
    'Random Forest': (RandomForestRegressor(), {'model__n_estimators': [10, 50, 100], 'model__max_depth': [None, 5, 10, 15]}),
    'XGBoost': (XGBRegressor(), {'model__n_estimators': [50, 100, 200], 'model__max_depth': [None, 5, 10, 15]})
}

# Perform hyperparameter tuning and evaluation for each model
for model_name, (model, param_grid) in models.items():
    pipeline = Pipeline([
        ('preprocessing', column_Transformer), # Include any necessary preprocessing steps
        ('model', model),
    ])

    # Create a GridSearchCV object
    grid_search = GridSearchCV(
        pipeline,
        param_grid,
        scoring='r2', # Use mean squared error as the scoring metric
        cv=5, # 5-fold cross-validation
        n_jobs=-1 # Use all available CPU cores
    )

    # Fit the GridSearchCV object on the training data
    grid_search.fit(X_train, y_train)

    # Print the best hyperparameters and corresponding mean squared error
    best_params = grid_search.best_params_
    best_mse = -grid_search.best_score_
    print(f"{model_name} - Best Hyperparameters: {best_params}, Best Mean Squared Error: {best_mse}")

    # Make predictions on the test set using the best model
    y_pred = grid_search.predict(X_test)

    # Evaluate the model on the test set
    r2 = r2_score(y_test, y_pred)

    # Append results to the DataFrame
```

```

results_df = results_df.append({
    'Model': model_name,
    'y_test': y_test.values,
    'y_pred': y_pred,
    'R2 Score': r2
}, ignore_index=True)

```

Decision Tree - Best Hyperparameters: {'model__max_depth': 10}, Best Mean Squared Error: -0.9206793331169172
Random Forest - Best Hyperparameters: {'model__max_depth': None, 'model__n_estimators': 100}, Best Mean Squared Error: -0.9461783055997636
XGBoost - Best Hyperparameters: {'model__max_depth': 5, 'model__n_estimators': 200}, Best Mean Squared Error: -0.95688975410082

```

In [57]: # Display the results DataFrame
results_df

```

Out[57]:


	Model	y_test	y_pred	R2 Score
0	Decision Tree	[1054454.4, 1591920.42, 1415473.91, 498580.87,...	[1085797.6744444445, 1095429.9233333333, 14192...	0.919111
1	Random Forest	[1054454.4, 1591920.42, 1415473.91, 498580.87,...	[1070848.9127000018, 1124118.5976999989, 13974...	0.942977
2	XGBoost	[1054454.4, 1591920.42, 1415473.91, 498580.87,...	[1053759.5, 1379877.9, 1402597.4, 820874.1, 17...	0.964359

```


In [58]: for i, row in results_df.iterrows():
plt.figure(figsize=(8, 6))
sns.regplot(x=row['y_test'], y=row['y_pred'], line_kws={'color': 'red', 'linestyle':
plt.title(f"{row['Model']} - Actual vs Predicted", fontdict={'fontsize': 20, 'color': 'G
plt.xlabel('Actual (y_test)', fontdict={'fontsize': 14, 'color': 'orange', 'fontweight': '
plt.ylabel('Predicted (y_pred)', fontdict={'fontsize': 14, 'color': 'orange', 'fontweight
plt.xticks(fontweight='bold', fontsize=12)
plt.yticks(fontweight='bold', fontsize=12)
plt.show()


```

Decision Tree - Actual vs Predicted



Random Forest - Actual vs Predicted





Insights:

XGBOOST performs the best

FORECAST the Weekly_Sales for each Stores for the NEXT 12 WEEKS.

Time Series for forecasting:

```
In [59]: #Apply time series analysis to identify seasonal trends:




# Use seasonal decomposition to identify seasonal, trend, and residual components
result = sm.tsa.seasonal_decompose(df_Walmart['Weekly_Sales'], model='additive', period=


# Plot the decomposed components
fig, ax = plt.subplots(4, 1, figsize=(20, 40))
result.observed.plot(ax=ax[0])
ax[0].set_ylabel('Observed')
result.trend.plot(ax=ax[1])
ax[1].set_ylabel('Trend')
result.seasonal.plot(ax=ax[2])
```



```
ax[2].set_ylabel('Seasonal')
result.resid.plot(ax=ax[3])
ax[3].set_ylabel('Residual')

plt.tight_layout()
plt.show()
```





In []:

Future Possibilities:

1. **Advanced Predictive Modeling:** While the current models (Decision Tree, Random Forest, XGBoost) provided satisfactory results, future work could explore more advanced forecasting models, including neural networks and time series models like ARIMA or SARIMA.
2. **Dynamic Feature Engineering:** Incorporate additional features or engineered features to enhance model accuracy. For example, adding promotional events, economic indicators, or regional data could provide a more comprehensive understanding of sales dynamics.
3. **Fine-Tuning Models:** Further hyperparameter tuning and optimization of the selected models could improve their predictive performance. Grid search techniques and Bayesian optimization can be explored for this purpose.
4. **Dynamic Inventory Management:** Utilize insights gained to optimize inventory management strategies.

Understanding the impact of external factors can help Walmart plan for demand fluctuations and reduce stockouts or overstock situations.

5. **Real-Time Data Integration:** Implement real-time data integration to capture the most recent information. This would enable Walmart to adapt its strategies promptly based on changing economic conditions, consumer behavior, or external events.
6. **Geospatial Analysis:** Incorporate geospatial analysis to understand the impact of location on sales. Factors such as population density, competition, and local events can play a crucial role in store performance.
7. **Customer Segmentation:** Analyze customer segmentation to tailor marketing and sales strategies for different consumer groups. This can be achieved by leveraging data on purchasing behavior, demographics, and preferences.
8. **Collaborative Filtering:** Implement collaborative filtering techniques to provide personalized recommendations to customers. This can enhance the customer shopping experience and potentially boost sales.
9. **Real-time Dashboard:** Developing a real-time dashboard to visualize the predicted versus actual sales, key performance indicators, and relevant metrics could assist stakeholders in making informed decisions promptly. This would enable quick reactions to emerging trends or anomalies.

By continuously refining models, incorporating new data sources, and adopting advanced analytics techniques, Walmart can stay ahead in the highly competitive retail landscape and make informed, data-driven decisions.

This concludes the analysis and provides a roadmap for future enhancements and strategic planning for Walmart's sales forecasting.

Conclusion:

In conclusion, the analysis of Walmart's sales data has provided valuable insights into various factors affecting weekly sales across its multiple stores. The key findings are:

1. **Unemployment Impact:** There is a negative correlation between Weekly Sales and Unemployment Rate. Some stores show a stronger negative correlation, indicating that economic conditions may influence sales.
2. **Seasonal Trends:** The analysis suggests the presence of seasonal trends, especially during holiday weeks, which significantly impact Weekly Sales. This insight can help Walmart better prepare for peak sales periods.
3. **Temperature Influence:** Temperature appears to have an impact on Weekly Sales, with a slight positive correlation. Further analysis indicates that during colder months, sales tend to increase.
4. **CPI and Fuel Prices:** The Consumer Price Index (CPI) shows varying impacts on Weekly Sales across different stores. Fuel prices also exhibit an influence on sales, but the relationship is not as straightforward.
5. **Store Performance:** The analysis identifies top-performing and worst-performing stores based on yearly sales trends. Understanding the performance of individual stores can guide strategic decision-making.
6. **Statistical Analysis:** Detailed statistical measures, including range, average sales, standard deviation, and coefficient of variance, were computed to assess the significance of the difference between the highest and lowest performing stores.
7. **The predictive modeling:** The predictive models, particularly XGBoost, demonstrated strong performance in forecasting weekly sales. However, the future possibilities outlined suggest avenues for further refinement and enhancement of the forecasting process.

References:

- Box, G. E. P., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). Time Series Analysis: Forecasting and Control. Wiley.
- Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '16).
- Hyndman, R. J., & Athanasopoulos, G. (2018). Forecasting: Principles and Practice. OTexts.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, 2825–2830.
- Zhang, G. P. (2003). Time series forecasting using a hybrid ARIMA and neural network model. Neurocomputing, 50, 159-175.